

Multi-head attention softens but does not remove the input bottleneck

